

Berlin JUG 2015

Motivations

- . New map-like API, functional + async
- . ConcurrentMap? not distributed in mind
 - @ show code: Infinispan Cache is ConcurrentMap, show put/replace previous value issues
- . Two ways to solve the issue: AdvancedCache.withFlags() or JCache standard API
- . But misses point: to better scale remote or persistence calls, we need to make things async!
- . JCache has no async call support
- . AdvancedCache has async call support but based on pre-Java8 Futures
- . Not want to just making async transport -> end result needed, e.g. did it fail?
- . AdvancedCache = bazar, so rather than change it, create a new API with async as top prio

New API

- . Infinispan 8 -> Java 8: so use better async support + functional constructs to define new API
- . One API -> Three map APIs: ReadOnly, WriteOnly, ReadWrite
 - @ code: construction (already prepared, so expand it)

Single Key

- . WriteOnlyMap & ReadOnlyMap
 - . WO operations do not read previous value
 - . WO single-key ops return CF<Void> since nothing can be derived from function
 - . RO operations do not acquire locks on key, WO operations do
 - . Returns of RO ops are flexible, could be anything from value to metadata
 - . Read view has a find() method to return an Optional of the value
 - . Read view offers a get() as shortcut for when we have certainty (throws NoSuchElementException)
 - . Lacks Haskell do notation or Scala fors but CompletableFuture's can be composed
 - @ code: WO eval K -> RO eval K get() -> WO eval K remove() -> RO eval K find()
- . ReadWriteMap

- . RW operations read previous value and acquire locks
- . RW can be used for ifAbsent operations, replaces, conditional replaces...etc
- . Conditions can be based on value or metadata information, e.g. version
- . Lifespan in example uses Java 8 API but could use concurrent TimeUnit too -> milliseconds
- . Unless the value is known, use Value-bearing eval functions to avoid capturing external vars
- . If not capturing, lambdas can be cached and makes clustering easier (more later)

@ code: RW eval K, V if absent w/ lifespan -> RO eval K lifespan (w/ sleep in between)

Multi Key

- . Pass in keys (+values) and function to execute on each key(+value), return Traversable
- . Traversable is a lazy & pull-style API
- . Exposes methods for working with returned data from each function
- . WO multi key operations still return CF<Void>
- . Normally Traversable order matches order of input collection (though not guaranteed)
- . Returning view useful for returning value + metadata information

@ code: WO evalMany KVs, RW evalMany Ks return previous view, Traversable filter + print

- . Also available:
 - . RO.keys() for traversing all keys, RO.entries() for traversing keys, values and metadata
 - . WO.evalAll() and RW.evalAll() for updating/removing all keys
- . ^ Traversable order not guaranteed
- . push-style API?
 - . in push-style, data pushed to callbacks provided by user instead of lazily fetching them
 - . this is the approach used by Rx, but we decided against it because:
 - . interested in providing Rx-like API but did not want hard dependencies in core
 - . didn't want to reinvent an Rx-like API -> not trivial -> backpressure, flow control
 - . pull-style API can be asynchronous since user can decide to use Traversable later

- . Traversable providers clear separation between inter & terminating ops
- . ^ which makes it a good abstraction to avoid unnecessary computation
- . strive to keep this separation clear results in no manual iteration methods
- . so, no iterator() nor spliterator() methods
- . associated with manual, user-end iteration, and we want to avoid such thing
- . majority of cases, Infinispan knows best how to exactly iterate over the data

Clustering & Listeners

- . Demonstrate how functional maps work in clustered environments
- . Live code will also use listeners to see effects in remote nodes
- . All listener events are post-event! Events received after event has happened
 - . vs Infinispan Cache events, which are pre & post
 - . only post because listener events are only there to find out what's happened
 - . pre-events can hint at the possibility of being able to alter execution
 - . for pre-operation work or to alter flow, use interceptors
- . Listener registration methods return AutoCloseable -> handler to deregister listeners
- . Start w/ write listeners: for write events in W0 or RW maps
 - . can be added passing a lambda to listeners.on() method, and gets a read view of entry write
 - . write events can't differentiate create/modifid events, since no previous value is available
 - . they can differentiate between remove and create/modified since value can be queried
- . W0 & RW ops: lambda is replicated to other nodes...
 - . So it needs to be somehow serializable!
- @ code: start with W0 map with write listener definition (both) -> W0 to store a constant value
 - . make sure listener is added in both nodes!
- @ code: show failure to serialize and quick and dirty workaround
 - . explain downsides of quick and dirty workaround => payloads are huge and slow
 - . show helper class with utility externalizer based

```
functions => quick and tiny payloads!  
@ code: show read-write listeners, RW.evalMany w/ metadata  
(setValueMetasIfAbsentReturnBoolean)  
  . add read-write listener in node1, onCreate + onModify  
  . do evalMany in node2 to create new entries  
  . do individual eval in node1 to modify one entry and see  
modified events
```